

# IT5100B

Industry Readiness  
*Stream Processing*

## LECTURE 3

Reactive Programming


**FOO** Yong Qi  
yongqi@nus.edu.sg

## CONTENTS

- What and Why Reactive Programming
- Project Reactor
- Flux/Mono Basics
- Additional Reactive Topics
- Reactive Programming Tips

# WHAT/WHY REACTIVE PROGRAMMING

# REACTIVE PROGRAMMING



**Declarative** programming paradigm concerned with **data streams** and **propagation of change**

# SMALL RESTAURANT

```
def chicken_chop():  
    chicken = grill_chicken()  
    fries = fry_fries()  
    return chicken + fries
```




Image source: Shutterstock

You are the chef taking and making orders

Suppose there is only one customer:

1. Receive the order
2. Grill the chicken
3. Fry the fries
4. Put chicken and fries together
5. Deliver order

# MEDIUM RESTAURANT




Suppose now you're the **head chef** and you have many chicken grillers and fries fryers

Suppose there is only one customer:

1. Receive the order
2. Tell the griller to grill chicken
3. Once you have the chicken, tell fryer to fry fries
4. Once you have the fries, put them together and deliver

# MEDIUM RESTAURANT




Suppose now you're the **head chef** and you have many chicken grillers and fries fryers

Suppose there is only one customer:

1. Receive the order
2. Tell the griller to grill chicken
3. Wait for chicken to come
4. Tell fryer to fry fries
5. Wait for fries to come
6. Put them together and deliver

# MEDIUM RESTAURANT




Suppose now you're the **head chef** and you have many chicken grillers and fries fryers

Suppose there is only one customer:

1. Receive the order
2. Tell the griller to grill chicken
3. Tell fryer to fry fries
4. Wait for both to come
5. Put them together and deliver

# BIG RESTAURANT




Suppose now business is good and you have many customers:

For each customer:

1. Receive the order
2. Tell the griller to grill chicken
3. Tell fryer to fry fries
4. Wait for both to come
5. Put them together and deliver



# BIG RESTAURANT



Suppose now business is good and you have many customers:

For each customer:

1. Receive the order
2. Tell the griller to grill chicken
3. Tell fryer to fry fries
4. When both comes:
  1. Put them together and deliver

# NON-BLOCKING

Life is all about **waiting**:

- Waiting for messages to come in
- Waiting for HTTP request to send response
- Waiting for the girl you're seeing to text you back

Instead of doing nothing while waiting, carry on with your life

When things come in, **react** to them

# NON-BLOCKING BEHAVIOUR

How do we introduce non-blocking behaviour into programs?

1. Create new classes that have a `whenSomethingComes(doSomething)` method
2. Re-write programs using this behavior
3. Write other features like schedulers to ensure performance is optimized, error handling, etc

Can we expect developers to do this?

# CALLBACKS

How to give your chef a chicken:

1. Look in refrigerator
2. If there are chickens, give them one
3. Otherwise, get one from the store and give it to them

Callback hell!

```
refrigerator.getChickens(new Callback<List<Chicken>>() {  
    public void onSuccess(List<Chicken> ls) {  
        if (ls.isEmpty()) {  
            store.buyChicken(new Callback<Chicken>() {  
                public void onSuccess(Chicken c) {  
                    chef.give(c);  
                }  
                public void onError(Throwable t) {  
                    chef.noMoreChickens();  
                }  
            });  
        } else {  
            chef.give(ls.get(0));  
        }  
    }  
    public void onError(Throwable t) {  
        store.buyChicken(new Callback<Chicken>() {  
            public void onSuccess(Chicken c) {  
                chef.give(c);  
            }  
            public void onError(Throwable t) {  
                chef.noMoreChickens();  
            }  
        });  
    }  
});
```

# COMPLETABLEFUTURES

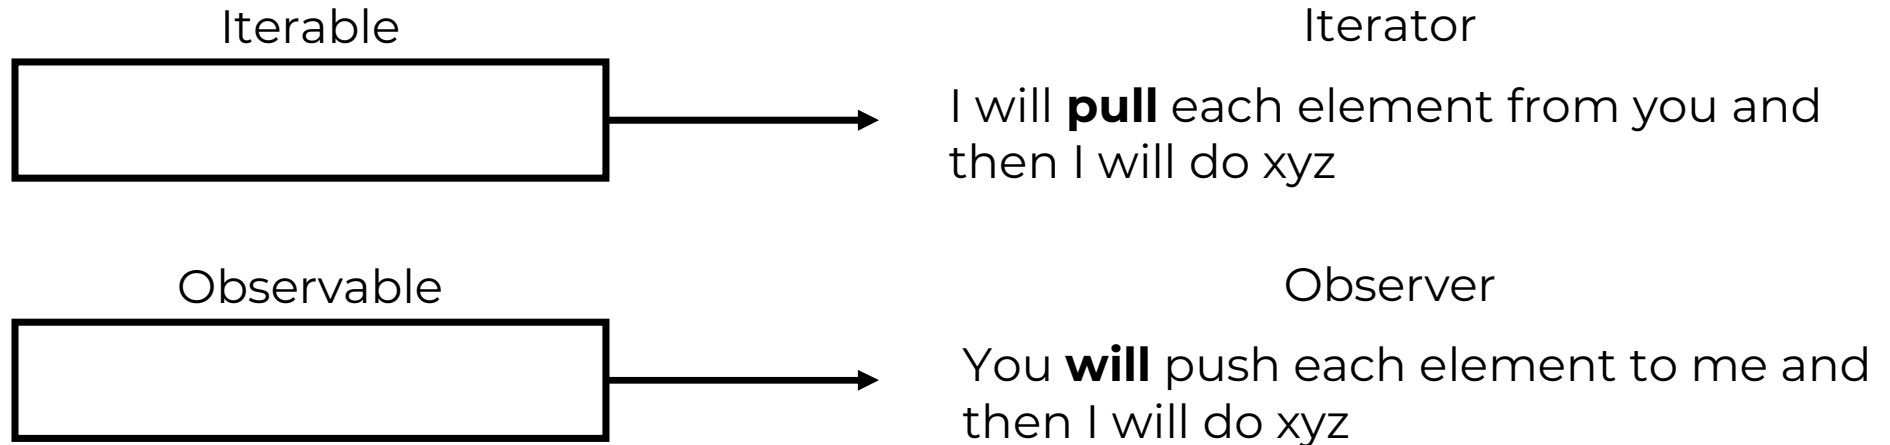
How to give your chef a chicken:

1. Look in refrigerator
  2. If there are chickens, give them one
  3. Otherwise, get one from the store and give it to them
- Slightly better with lambda expressions, still hard to read!
  - Clunky error handling!
  - What about non-blocking streams?

```
refrigerator.getChickens().handle((ls, exp) -> {  
    if (ls != null && ls.isEmpty()) {  
        store.buyChicken().handle((chicken, exp) -> {  
            if (chicken != null) {  
                chef.give(chicken);  
            } else {  
                chef.noMoreChickens();  
            }  
        });  
    } else if (ls != null) {  
        chef.give(ls.get(0));  
    } else {  
        store.buyChicken().handle((chicken, exp) -> {  
            if (chicken != null) {  
                chef.give(chicken);  
            } else {  
                chef.noMoreChickens();  
            }  
        });  
    }  
});
```

# ITERATOR VS OBSERVER

## A PARADIGM SHIFT



Aside from underlying implementation, the way you will interact with iterables (like lists) and observables (reactive streams) is the same!

# ITERATOR VS OBSERVER

## A PARADIGM SHIFT

Iterator code

```
keyPressList.filter(x -> x != null)  
  .map(Event::withId)  
  .forEach(System.out::println);
```

Observer code

```
keyPressObservable.filter(x -> x != null)  
  .map(Event::withId)  
  .forEach(System.out::println);
```

The code we write for two completely different models should look virtually the same!

# NON-BLOCKING REACTIVE PROGRAMMING

```
refrigerator.getChickens()           // get list of chickens
  .next()                           // just take one
  .doOnError(chef::somethingHappened) // tell chef if error occurred
  .onErrorComplete()                // recover from error if occurred
  .switchIfEmpty(store.getChicken()) // if nothing from fridge, get from store
  .singleOptional()                  // potentially got nothing from whole process
  .subscribe(oc ->
    oc.ifPresentOrElse(chef::give, // if chicken received, give to chef
      chef::noMoreChickens),        // if no chickens received, tell chef
    chef::somethingHappened);        // if error occurred at the end, tell chef
```

Much cleaner way to express non-blocking code in familiar style  
with robust error handling



# REACTIVE PROGRAMMING IN THE WILD

## CASE STUDY: TOMCAT VS NETTY

	Apache Tomcat	Netty (Reactive)
Throughput	2310 req/s	3421 req/s
Memory	521MB	62MB
CPU Usage	~70%	~50%
Threads	220	19

“[Netty can] service **50,000+** messages per second from approximately **30,000 connected clients** on a **commodity Intel server** costing approximately \$850”

- Matt Rhodes, RhoMoSoft, February **2009**

<https://medium.com/@skhatri.dev/springboot-performance-testing-various-embedded-web-servers-7d460bbfdb1b>

Maximizing thread usage by non-blocking improves performance and allows applications to handle more concurrent requests

# KEY POINT #1

**Reactive programming** frameworks allow us to write **non-blocking** programs in a familiar **declarative** style

# PROJECT REACTOR

# REACTIVE PROGRAMMING LIBRARIES

## PROJECT REACTOR

**Fully non-blocking I/O** with **functional API** over **typed [0|1|N] sequences**

Works nicely with popular frameworks, drivers and protocols:

- Spring Boot and WebFlux
- CloudFoundry Java Client
- Reactive Relational Database Connectivity (R2DBC)

Project Reactor is an alternative to RxJava

# PROJECT REACTOR

## MAIN CLASSES

`Flux<T>`

0 or N elements

`Mono<T>`

0 or 1 element

Non-reactive equivalents:

`Stream<T>`


`Optional<T>`

---

See the [Flux](#) and [Mono](#) API documentations

# ASSEMBLY LINE

## MENTAL MODEL OF FLUX AND MONO




Each Flux or Mono object is a Publisher that publishes new data

Think of this as a conveyor belt carrying items

# ASSEMBLY LINE


MENTAL MODEL OF FLUX AND MONO



Each operation on a Flux or Mono is like a workstation that creates a new Flux or Mono

# ASSEMBLY LINE

MENTAL MODEL OF FLUX AND MONO




Each Flux or Mono can be **subscribed** to, which produces subscriptions known as Disposables

Generally, nothing happens until a subscription occurs (lazy!)



# MARBLE DIAGRAMS


MENTAL MODEL OF FLUX AND MONO



Visual representations of Flux and Mono and signals produced through the flow of time

# MARBLE DIAGRAMS

## SIGNALS




One of three signals can be emitted from a Flux or Mono:

1. An object (next signal)
2. A complete signal (no more elements after)
3. An error signal (no more elements after-ish\*)

---

\*Flux#onErrorContinue can cause upstream publishers to recover from errors but this should be rarely used

# FLUX VS MONO




Flux produces 0 or more elements, Mono produces at most one

Both Flux and Mono are Publishers in the Reactive Streams specification

# MARBLE DIAGRAMS


## OPERATORS



Instance operators operate on Flux/Mono objects and produce a new Flux/Mono

# MARBLE DIAGRAMS


## STATIC OPERATORS



Static operators return new Flux/Mono may also receive publishers

# MARBLE DIAGRAMS


## UPSTREAM SIGNALS



Signals can be sent upstream to earlier publishers  
subscribe signals to earlier publishers to beginning emitting events

# MARBLE DIAGRAMS

## SIDE EFFECTS



Side effects can be executed on different signals; some happen before (left) or after (right) the signal is sent

# FLUX/MONO BASICS



# PUBLISHER CREATION

## JUST/FROM

```
Mono<Integer> m1 = Mono.just(1);  
Flux<String> f2 = Flux.just("hello", "world!");  
Mono<Object> m3 = Mono.justOrEmpty(null);
```

The just series of static methods wrap element(s) in a Flux/Mono

```
Mono<Integer> m4 = Mono.fromSupplier(() -> 1);  
Flux<Double> f5 = Flux.fromStream(Stream.of(1, 1.5, 2));
```

The from series of static methods create Flux/Mono from different data structures

# PUBLISHER CREATION

## FROM EXTERNAL SERVICES

```
Flux<User> m = db.sql("SELECT * FROM users")  
    .map(User::fromDb)  
    .all();
```


Generally, reactive data sources are not created programmatically:

- R2DBC SQL query to a database gives Flux/Mono
- ReactorKafka gives Flux/Mono on event production to Kafka message broker
- Etc.

# PUBLISHER CREATION

## DEFER

```
AtomicInteger i = new AtomicInteger(1);  
Mono<Integer> m = Mono.defer(() -> {  
    int x = i.getAndIncrement();  
    System.out.println("Mono #" + x);  
    return Mono.just(x);  
});
```



Lazily supply a Publisher every time a subscription is made; each subscriber will get a subscriber-specific instance

# PUBLISHER CREATION

## GENERATE

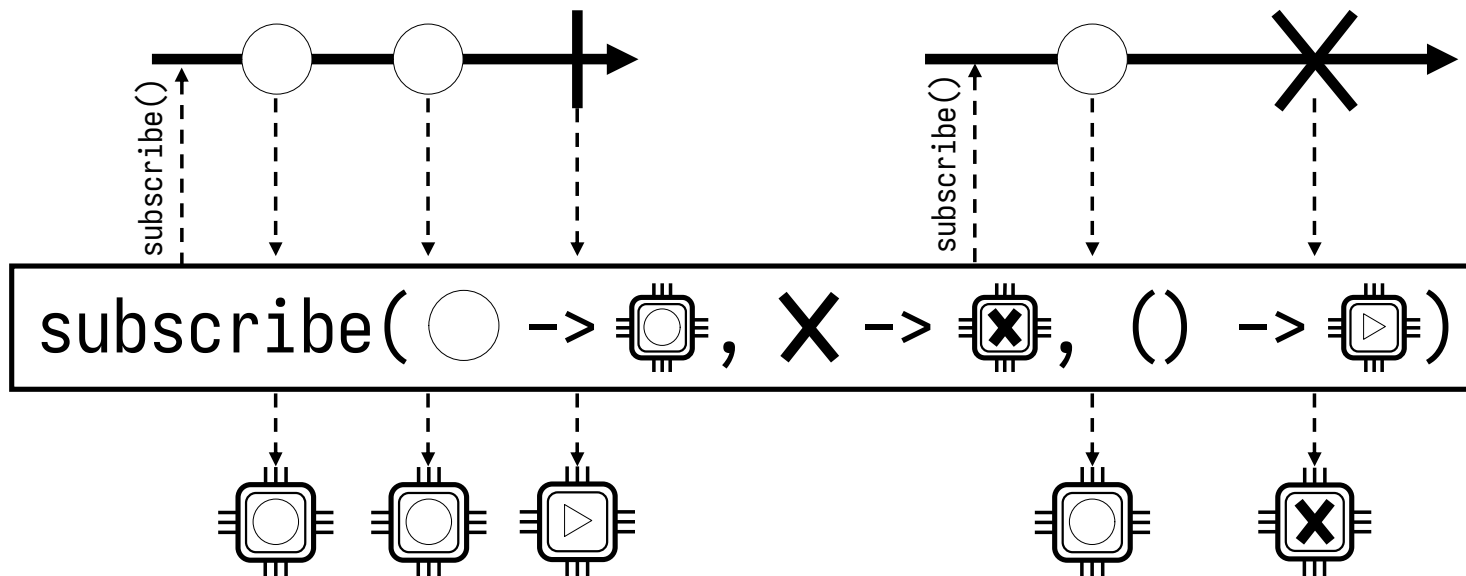
```
Flux<Integer> f = Flux.generate(() -> 1, (i, s) -> {  
    s.next(i); // push next element into sink  
    if (i == 10) s.complete(); // send complete signal  
    return i + 1; // return next state  
});
```

Stateful synchronous generation of elements (provides a lot of fine-grain control)

Similar to `Stream#iterate`

# SUBSCRIPTION

## SUBSCRIBE VARIANTS



Triggers the publisher to start emitting signals; each signal can be handled by a different consumer; subscribe returns a Disposable which can be disposed (cancels the subscription)

# PUBLISHER TRANSFORMS

## MAP AND FILTER

```
Flux.range(0, 10)  
    .map(x -> x * 2)  
    .filter(x -> x % 3 == 0)  
    .subscribe(System.out::println);
```

Simple transforms can be done with `map` and `filter` as per usual

# EXERCISE

## FLUX CREATION

Write a static method `cycle` that receives a `List` and produces a `Flux` that cycles over the elements repeatedly (potentially unbounded). Use `generate` (and `map` too).

```
Flux<String> f = cycle(List.of("a", "b", "c"));  
f.take(5).subscribe(System.out::println); // a, b, c, a, b
```

# SOLUTION

## CYCLE

```
public static <T> Flux<T> cycle(List<? extends T> ls) {  
    if (ls.isEmpty()) return Flux.empty();  
    return Flux.<Integer, Integer>generate(  
        () -> 0, // index of first item  
        (i, s) -> {  
            s.next(i); // push index  
            return (i + 1) % ls.size(); // next index  
        })  
        .map(ls::get); // map indices to elements  
}
```



# PUBLISHER TRANSFORMS

## FLATMAP

`flatMap` works as you would expect, except with (mainly) three variants:

	<code>flatMap</code>	<code>flatMapSequential</code>	<code>concatMap</code>
Subscription of inner publishers	Eager	Eager	Only after previous inner fluxes have completed
Ordering of signals	Next signal is first emission from any inner flux	Ordered like concatenation	Ordered (concatenation)
Interleaving	Yes	No	No (concatenation)

# PUBLISHER TRANSFORMS


## GENERAL TRANSFORMS

```
Function<Flux<Integer>, Flux<String>> f = s -> s.map(x -> x * 2)
    .filter(x -> x % 3 == 0)
    .map(x -> x + " haha");

Flux.range(1, 10)
    .transform(f)
    .subscribe(System.out::println);
}
```

For convenience and readability, you may use the general transform method that performs a bunch of transforms on a publisher; there are transform variants you may find useful too

# HOOKS



The `do` series of methods add side-effects to signals

# EXERCISE

## TRANSFORMS AS “AND THEN” OPERATIONS

```
public static Mono<Integer> getFromDb(int i) {  
    return Mono.just(i)  
        .delayElement(Duration.ofSeconds(1))  
        .doOnNext(x -> System.out.println("retrieved " + x));  
}  
public static Mono<Integer> putInDb(int i) {  
    return Mono.just(i)  
        .delayElement(Duration.ofSeconds(1))  
        .doOnNext(x -> System.out.println("put "+ x));  
}
```

Write a static method `increment` that receives an integer `i`, then gets it from the database, increases it by one, and puts it back in the database

# SOLUTION

TRANSFORMS AS “AND THEN” OPERATIONS

```
public static Mono<Integer> increment(int i) {  
    return getFromDb(i).map(x -> x + 1)  
        .flatMap(Main::putInDb);  
}
```

# MERGES

There are many ways to merge publishers together:

- `concat`: concatenate publishers
- `merge`
- `zip`: combine n-wise to n-tuples
- `and/when`: coordinate termination
- `firstWith...:` Choose first emitter
- `...ifEmpty`: Reactive if-else
- `then`: Chain publishers

# ERROR HANDLING

There are many error handling operations:

- `onError...`: perform some transform when error signal received
- `doOnError`: perform side-effect on error signal
- `retry`: re-subscribe to the publisher if given error signal
- `subscribe`: subscribe with custom error handler

# EXERCISE

## MERGES AND ERROR HANDLING

```
public static Mono<Integer> getRandInt() {  
    return Mono.just(rng.nextInt(10))  
        .doOnNext(x -> System.out.println("got " + x));  
}
```

Define a method `divide` that gets two integers `a` and `b` reactively and returns `a / b` reactively. If there is a `ArithmeticException`, print an error message and complete with no elements



# SOLUTION

## MERGES AND ERROR HANDLING

```
public static Mono<Integer> divide() {  
    return getRandInt().zipWith(getRandInt())  
        .map(x -> x.getT1() / x.getT2())  
        .doOnError(ArithmeticException.class, System.out::println)  
        .onErrorComplete(ArithmeticException.class);  
}
```

# ORIGINAL EXAMPLE WALKTHROUGH

```
refrigerator.getChickens()           // get list of chickens
  .next()                           // just take one
  .doOnError(chef::somethingHappened) // tell chef if error occurred
  .onErrorComplete()                 // recover from error if occurred
  .switchIfEmpty(store.getChicken()) // if nothing from fridge, get from store
  .singleOptional()                  // potentially got nothing from whole process
  .subscribe(oc ->
    oc.ifPresentOrElse(chef::give, // if chicken received, give to chef
      chef::noMoreChickens),        // if no chickens received, tell chef
    chef::somethingHappened);        // if error occurred at the end, tell chef
```

# ADDITIONAL REACTIVE TOPICS

# CONTROLLING BACKPRESSURE

## PUSH/PULL MODEL

```
myFlux().subscribe(new BaseSubscriber<>() {  
    protected void hookOnNext(Integer value) {  
        System.out.println(value);  
        request(1);  
    }  
    protected void hookOnComplete() {  
        System.out.println("no more!");  
    }  
});
```

Create custom subscribers that pull available data on demand,  
controlling backpressure

# HOT/COLD PUBLISHERS

So far most publishers we have seen are cold:

Cold publisher—nothing happens until subscribe, each subscribe triggers publication (most assembly operations)

Hot publisher—publish data eagerly, when new subscribers only see signals after subscription (unless cached) (just)

# THREADING AND SCHEDULERS

## HANDLING CONCURRENCY AND SCHEDULERS

```
return Mono.fromSupplier(() -> {  
    try { Thread.sleep(1000); } catch (Exception e) { }  
    return 1;  
}).publishOn(Schedulers.boundedElastic());
```

Reactor is **concurrency agnostic**, you get to choose how threading is done

# MANY MORE!

Many more topics on reactive programming with Reactor can be found on their [amazing reference documentation](#)

# REACTIVE PROGRAMMING TIPS



# REACTIVE ALL THE WAY!

In Streams/Optionals, eventually we will reduce/collect/get.  
When do we block to retrieve elements from publishers?

Ideally, **never!**

Let the framework decide when to block; if you block, you wait!

# DON'T BREAK THE CHAIN

```
public static Mono<Mono<Void>> something() {  
    return Mono.just(1)  
        .map(x -> doStuff());  
}  
public static Mono<Void> doStuff() {  
    return Mono.just(1)  
        .doOnNext(x -> System.out.println("stuff"))  
        .then();  
}
```

```
public static Mono<Void> something() {  
    return Mono.just(1)  
        .flatMap(x -> doStuff());  
}  
public static Mono<Void> doStuff() {  
    return Mono.just(1)  
        .doOnNext(x -> System.out.println("stuff"))  
        .then();  
}
```

Cold publishers are never triggered until subscribed; subscribing to a publisher will not trigger subscription unconnected ones

## **KEY POINT #2**

Stay reactive!

## CONTENTS

- What and Why Reactive Programming
- Project Reactor
- Flux/Mono Basics
- Additional Reactive Topics
- Reactive Programming Tips

## KEY POINTS

- **Reactive programming** frameworks allow us to write **non-blocking** programs in a familiar **declarative** style
- Stay reactive!